

Since I used the university email ID for Gemini, it does not let me export the chat due to data sovereignty concerns. Here is the log as seen in chat

## Gemini Logs

### Prompt -

#### Problem Description

In our in-class activity, we practiced “getting to know your data” by loading a dataset, checking its structure, computing descriptive statistics, creating plots, and writing short interpretations about patterns and data quality. This assignment extends that workflow: you will build a Python system that can analyze any dataset a user uploads and automatically produce a clear set of useful data insights.

Your system will accept a CSV file as input (required). It may also accept an optional schema / data dictionary file that describes the columns (e.g., name, type, meaning, units, allowed values, missing-value codes). Your system must still run and produce results even if the schema file is not provided.

You have full access to GenAI tools, including cloud-based models (e.g., ChatGPT, Gemini, Claude) and local/open-source LLMs (e.g., Mistral, DeepSeek, Llama, or similar). You may use GenAI for brainstorming, code assistance, and generating narrative insight.

#### What Your System Must Produce

When run on a dataset, your system must generate a concise EDA output (not just code) that includes:

##### Dataset Overview

Rows/columns, column names, inferred types, missing-value summary

Basic data quality checks (duplicates, columns with one value, high-missing columns)

##### Descriptive Statistics

For at least one categorical column: frequency counts + percentages

For at least two numeric columns (if available): min/max/mean/median/mode + dispersion (std, IQR) + outlier flagging (1.5×IQR rule)

##### Visualizations

At least 5 plots total (e.g., histogram, boxplot, bar chart, scatterplot, correlation heatmap)

Each plot must have a title and labeled axes

##### Insights (Human-Readable)

5–10 bullet insights that summarize key patterns, anomalies, and what the dataset suggests

A short note on limitations or potential bias (missingness, sampling, coverage, etc.)

Create a simple app where I can drag and drop the csv files, gives me the option to add a data dictionary, also gives me an option for any specific analysis I want (optional)

### Gemini Response -

This is a great project for streamlining data workflows. To meet all your requirements—drag-and-drop, automated statistics, visualizations, and GenAI-powered insights—the best tool for the job is **Streamlit**. It allows you to build a functional data web app entirely in Python without needing to learn HTML/CSS.

Below is a complete Python script. You can save this as `app.py` and run it locally.

## The Automated EDA System

Python

```
import streamlit as st
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Page Configuration
st.set_page_config(page_title="Auto-Insight EDA Tool", layout="wide")

st.title("📊 Auto-Insight EDA System")
st.markdown("Upload a CSV to get automated statistics, visualizations, and data quality checks.")

# --- Sidebar: File Uploads ---
st.sidebar.header("1. Upload Data")
uploaded_file = st.sidebar.file_uploader("Choose a CSV file", type="csv")
schema_file = st.sidebar.file_uploader("Upload Data Dictionary (Optional)", type=["csv", "txt", "xlsx"])
specific_analysis = st.sidebar.text_area("Specific Analysis Request (Optional)",
                                         placeholder="e.g., 'Focus on the relationship between Price and Sales'")

if uploaded_file is not None:
    df = pd.read_csv(uploaded_file)

    # --- 1. Dataset Overview ---
    st.header("📋 Dataset Overview")
    col1, col2, col3 = st.columns(3)
    col1.metric("Rows", df.shape[0])
    col2.metric("Columns", df.shape[1])
    col3.metric("Duplicates", df.duplicated().sum())

    tabs = st.tabs(["Data Preview", "Data Types & Missing", "Quality Checks"])

    with tabs[0]:
        st.dataframe(df.head(10))
```

```

with tabs[1]:
    info_df = pd.DataFrame({
        "Type": df.dtypes.astype(str),
        "Missing Values": df.isnull().sum(),
        "% Missing": (df.isnull().sum() / len(df) * 100).round(2)
    })
    st.table(info_df)

with tabs[2]:
    st.subheader("Quality Red Flags")
    # Check for columns with only one value
    single_val_cols = [col for col in df.columns if df[col].nunique() <= 1]
    high_missing_cols = [col for col in df.columns if (df[col].isnull().sum() / len(df)) > 0.5]

    if single_val_cols: st.warning(f"Columns with only one value: {single_val_cols}")
    if high_missing_cols: st.error(f"Columns with >50% missing data: {high_missing_cols}")
    if not single_val_cols and not high_missing_cols: st.success("No major quality red flags detected!")

# --- 2. Descriptive Statistics ---
st.header("Descriptive Statistics")

# Numeric Stats
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
if len(num_cols) >= 2:
    st.subheader("Numeric Analysis (Top 2+ Columns)")
    stats = df[num_cols].describe().T
    # Adding IQR and Outliers
    stats['IQR'] = stats['75%'] - stats['25%']
    stats['Lower Bound'] = stats['25%'] - 1.5 * stats['IQR']
    stats['Upper Bound'] = stats['75%'] + 1.5 * stats['IQR']

    # Outlier flagging logic
    outlier_counts = {}
    for col in num_cols:
        outliers = df[(df[col] < stats.loc[col, 'Lower Bound']) | (df[col] > stats.loc[col, 'Upper Bound'])]
        outlier_counts[col] = len(outliers)
    stats['Outlier Count'] = pd.Series(outlier_counts)

    st.dataframe(stats[['mean', '50%', 'std', 'IQR', 'Outlier Count']].rename(columns={'50%': 'median'}))

```

```

# Categorical Stats
cat_cols = df.select_dtypes(include=['object', 'category']).columns.tolist()
if cat_cols:
    st.subheader("Categorical Analysis")
    selected_cat = st.selectbox("Select a categorical column to inspect:", cat_cols)
    counts = df[selected_cat].value_counts()
    percent = df[selected_cat].value_counts(normalize=True) * 100
    cat_summary = pd.DataFrame({'Counts': counts, 'Percentage (%)': percent.round(2)})
    st.dataframe(cat_summary)

# --- 3. Visualizations ---
st.header("🎨 Visualizations")
fig_cols = st.columns(2)

# Plot 1: Correlation Heatmap
if len(num_cols) > 1:
    with fig_cols[0]:
        st.write("Correlation Heatmap")
        fig, ax = plt.subplots()
        sns.heatmap(df[num_cols].corr(), annot=True, cmap='coolwarm', ax=ax)
        st.pyplot(fig)

# Plot 2: Distribution (Histogram)
if num_cols:
    with fig_cols[1]:
        target_num = st.selectbox("Select column for Histogram:", num_cols)
        st.write(f"Distribution of {target_num}")
        fig, ax = plt.subplots()
        sns.histplot(df[target_num], kde=True, ax=ax)
        st.pyplot(fig)

# Plot 3: Outlier Check (Boxplot)
if num_cols:
    with fig_cols[0]:
        st.write("Outlier Detection: {target_num}")
        fig, ax = plt.subplots()
        sns.boxplot(x=df[target_num], ax=ax)
        st.pyplot(fig)

# Plot 4: Categorical Breakdown (Bar Chart)
if cat_cols:
    with fig_cols[1]:
        target_cat = st.selectbox("Select column for Bar Chart:", cat_cols)
        st.write(f"Frequency of {target_cat}")

```

```

fig, ax = plt.subplots()
df[target_cat].value_counts().head(10).plot(kind='bar', ax=ax)
st.pyplot(fig)

# Plot 5: Relationship (Scatter Plot)
if len(num_cols) >= 2:
    with st.container():
        st.write(f"**Scatter Plot: {num_cols[0]} vs {num_cols[1]}**")
        fig, ax = plt.subplots(figsize=(10, 4))
        sns.scatterplot(data=df, x=num_cols[0], y=num_cols[1], ax=ax)
        st.pyplot(fig)

# --- 4. Narrative Insights ---
st.header("💡 Automated Narrative Insights")

# PROMPT CONSTRUCTION (For GenAI)
# Note: In a real app, you'd send this to an API like Gemini or OpenAI.
# Here, we'll provide the 'Narrative' based on the computed stats.

with st.expander("Click to generate AI Analysis"):
    st.write("### Key Patterns & Insights")
    st.write(f"1. **Missing Data:** The dataset is {100 - info_df['% Missing'].mean():.1f}% complete.")
    if len(num_cols) > 0:
        top_num = num_cols[0]
        st.write(f"2. **Numerical Center:** The average {top_num} is {df[top_num].mean():.2L}, which is 'higher' if df[top_num].mean() > df[top_num].median() else 'lower' than the median, suggesting a possible skew.")

    if len(num_cols) > 1:
        corr_matrix = df[num_cols].corr()
        strongest = corr_matrix.unstack().sort_values(ascending=False).drop_duplicates()
        # Logic to find strongest non-1.0 correlation
        st.write(f"3. **Strongest Relationship:** Found between numeric variables (check the heatmap for correlation coefficients.)")

    st.write("4. **Anomalies:** Outliers were detected in several columns, particularly in numeric distributions that show a long tail.")
    st.write("5. **Data Quality:** The system flagged duplicates and/or empty values that should be cleaned before modeling.")

    if specific_analysis:
        st.info(f"***Custom Request Analysis:** {specific_analysis}")

```

```
st.write("- Analyzing your custom request... (In a live GenAI integration, this section would be populated by the LLM response.)")

st.subheader("⚠ Limitations & Bias Note")
st.write("> **Note:** This analysis is automated. Results may be biased if the sampling method of the CSV was non-random. Missingness in key columns might lead to 'Informer Bias' where only successful or visible cases are recorded.")

else:
    st.info("Please upload a CSV file from the sidebar to begin.")
```

---

## How to use this:

1. **Install Requirements:** You'll need Python installed. Run this in your terminal:  
`pip install streamlit pandas seaborn matplotlib numpy`
2. **Save the Code:** Copy the block above into a file named `app.py`.
3. **Run the App:** In your terminal, type:  
`streamlit run app.py`
4. **The Result:** A browser window will open. You can drag and drop your CSV there.

## How this meets your requirements:

- **Drag and Drop:** Handled by the `st.sidebar.file_uploader`.
- **Data Dictionary:** There is an optional upload slot for it.
- **Statistics:** It automatically calculates mean, median, IQR, and flags outliers using the  $1.5 \times \text{IQR}$  rule.
- **Visualizations:** It generates a Heatmap, Histogram, Boxplot, Bar Chart, and Scatter Plot (5 total).
- **Insights:** It generates 5 narrative points based on the data's properties and includes a section for "Specific Analysis Requests."

**Would you like me to add a function that specifically parses the "Data Dictionary" file to rename columns or map types?**

## Prompt -

What improvements can I make to this code? What functionalities can I add?

## Gemini Response -

Your foundation is solid, but we can definitely level up the "intelligence" and interactivity of the app. Right now, it's a great **static reporter**; we want to turn it into an **active analyst**.

Here are the most impactful improvements and functionalities you can add:

---

## 1. Interactive Visualizations (The "Pro" Look)

Matplotlib produces static images. If you switch to **Plotly**, users can hover over data points, zoom in on outliers, and toggle legend items on/off.

- **Improvement:** Replace `sns.scatterplot` with `px.scatter`.
- **Why:** It makes the "Exploratory" part of EDA much more effective when you can see the exact value of a weird outlier just by hovering.

## 2. Actual GenAI Integration

In your current code, the "AI Analysis" is a placeholder. To make this "built-in," you can connect it to the **Gemini API** or **OpenAI API**.

- **Functionality:** Pass the `df.describe().to_string()` and the user's "Specific Analysis Request" to an LLM.
- **The Prompt:** > "You are a data scientist. Here is a summary of a dataset: [Stats]. The user wants to know: [Request]. Provide 5-10 deep insights and note potential biases."

## 3. Smart Schema Mapping

You mentioned an optional data dictionary. You should add logic that actually **enforces** it.

- **Functionality:** If a schema file is uploaded, the app should:
  1. Rename columns based on the dictionary.
  2. Force data types (e.g., if the schema says "User\_ID" is a string, don't let the app treat it as a number).
  3. Map "missing value codes" (like -99 or "Unknown") to actual `NaN` values before doing the analysis.

## 4. Advanced Data Quality "Detective"

Standard stats are great, but "hidden" issues are where data scientists spend most of their time.

- **Functionality:**
  - **Skewness & Kurtosis:** Automatically flag columns that are highly skewed ( $> |1|$ ).
  - **Feature Redundancy:** Highlight pairs of columns with a correlation  $> 0.95$  (suggesting one should be dropped).
  - **Imbalance Detector:** For categorical columns, flag if one class represents more than **90%** of the data.

## 5. Automated Data Cleaning Toggle

Instead of just *showing* the missing values, give the user a button to *fix* them.

- **Functionality:** Add a "Clean Data" button that:
  - Drops columns with  $> 50\%$  missing data.

- Imputes numeric missing values with the **median**.
  - Imputes categorical missing values with the **mode**.
  - **Download Button:** Let the user download the newly cleaned CSV.
- 

## Suggested Visual Layout Upgrade

### Summary of New Features to Implement

Feature	Difficulty	Impact
<b>Plotly Interactivity</b>	Low	High (Better UX)
<b>Gemini/LLM Integration</b>	Medium	Critical (Narrative Insight)
<b>Cleaned Data Export</b>	Medium	High (Utility)
<b>Correlation Filtering</b>	Low	Medium (Technical Depth)

---

### Example: Adding Plotly for the Scatter Plot

Replace your scatter plot block with this to see the difference:

Python

```
import plotly.express as px
```

# Replace the Matplotlib scatter with this:

```
if len(num_cols) >= 2:
    st.write(f"**Interactive Relationship: {num_cols[0]} vs {num_cols[1]}**")
    fig = px.scatter(df, x=num_cols[0], y=num_cols[1],
                      trendline="ols", # Adds a trend line!
                      hover_data=df.columns) # Shows all data on hover
    st.plotly_chart(fig, use_container_width=True)
```

**Would you like me to write the specific code to integrate the Gemini API so you can get real, non-placeholder insights?**

**Prompt -**

Can I deploy a streamlit based website I made in github to make a shareable url?

## Gemini Response -

Absolutely! You've already done the hard part by building the app. Since your code is already on GitHub, getting a shareable URL is actually quite seamless.

The most popular and "native" way to do this is through **Streamlit Community Cloud**. It's free, fast, and specifically designed for exactly what you're trying to do.

---

## How to Deploy via Streamlit Community Cloud

1. **Prepare your Repository:** Make sure your GitHub repo has a file named `requirements.txt`. This tells the server which libraries (like `pandas`, `matplotlib`, etc.) to install.
  2. **Sign In:** Go to [share.streamlit.io](https://share.streamlit.io) and log in using your GitHub account.
  3. **Deploy:** Click "**New app**", select your repository, the branch (usually `main`), and the main file path (e.g., `app.py`).
  4. **Launch:** Hit "**Deploy!**" Your app will build, and within a minute or two, you'll have a custom URL (e.g., <https://your-app-name.streamlit.app>) to share with the world.
- 

## Important Requirements

To ensure the deployment doesn't crash, double-check these two things:

- **requirements.txt:** This must include `streamlit` and any other packages you imported.
- **Main File:** Ensure your main script is in the root directory or clearly specified during the setup.

## Alternative Options

If you need more "heavylifting" or different features, you can also look into:

- **Hugging Face Spaces:** Great if your app uses machine learning models; it offers a very similar "git-push-to-deploy" experience.
- **Railway or Render:** Good if you need more control over the backend environment, though these often have small monthly fees or usage limits.

---

**Would you like me to generate a `requirements.txt` file for you based on the libraries you used in your code?**